

Elasticity Model Comparison & Rolling Validation

This notebook evaluates multiple regression models for estimating price elasticity and predicting sales demand. We compare linear models, regularized models, robust regression, kernel-based Vector Regression, ensemble methods (Random Forest and Gradient Boosting), and Neural Net (MLP Regressor), both with and without time-decay weights.

Backtest Methodology

- **Training Window:** 495 weeks (rolling window of the most recent 495 weeks of data).
- **Test Target:** The next week immediately following the training window.
- **Rolling Step:** 1 week (we roll the window forward, re-train on the new 495-week sliding window, and predict the next week).
- **Key Output Metrics:** MAE, MAPE, RMSE, average estimated Elasticity, and parameter volatility.

Model Ranking Weights

- **Prediction Accuracy (MAPE):** 40% weight (lower MAPE is better)
- **Elasticity Plausibility:** 30% weight (we want elasticity between -0.5 and -1.5)
- **Elasticity Stability (Std):** 30% weight (lower parameter volatility is better)

```
In [1]: import numpy as np
import pandas as pd
import statsmodels.api as sm
from sklearn.linear_model import Ridge, HuberRegressor
from sklearn.svm import SVR
from sklearn.ensemble import RandomForestRegressor, GradientBoostingRegressor
from sklearn.neural_network import MLPRegressor
import matplotlib.pyplot as plt
import warnings
warnings.filterwarnings('ignore')
```

```
In [2]: # =====
# CONFIGURATION VARIABLE
# =====
RUN_DECAY_ONLY = False # Set to True to evaluate only decay-based models
# =====

# Load the weekly aggregated dataset
df = pd.read_csv('sales_demand.csv')
df['date'] = pd.to_datetime(df['date'])
df = df.sort_values('date').reset_index(drop=True)
print(f"Loaded {len(df)} weeks of sales demand data. RUN_DECAY_ONLY = {RUN_DECAY_ONLY}")
df.head()
```

Loaded 495 weeks of sales demand data. RUN_DECAY_ONLY = False

Out [2]:

		date	sku_id	unit_price	units_sold	cost_
-	-		IND_BASMATI_	.	.	.
-	-		IND_BASMATI_	.	.	.
-	-		IND_BASMATI_	.	.	.
-	-		IND_BASMATI_	.	.	.
-	-		IND_BASMATI_	.	.	.

Model Definition Helper

We implement a generic helper function to run each model type on a training slice.

```
In [3]: def run_model(model_type, train_slice, test_price, use_decay=True, decay_
n_obs = len(train_slice)
if use_decay:
    max_date = train_slice['date'].max()
    t_diff_weeks = (max_date - train_slice['date']).dt.days / 7.0
    weights = decay_rate ** t_diff_weeks
else:
    weights = np.ones(n_obs)

log_p = np.log(train_slice['unit_price'].values).reshape(-1, 1)
log_q = np.log(train_slice['units_sold'].values)
log_p_base = np.log(test_price)
h = 1e-4

if model_type == 'Linear':
    X = sm.add_constant(log_p)
    model = sm.WLS(log_q, X, weights=weights).fit()
    elasticity = model.params[1] if len(model.params) > 1 else -1.5
    pred_log_q = model.params[0] + elasticity * log_p_base
    pred_q = np.exp(pred_log_q)
    return elasticity, pred_q

elif model_type == 'Ridge':
    model = Ridge(alpha=1.0)
    model.fit(log_p, log_q, sample_weight=weights)
    elasticity = model.coef_[0]
    pred_log_q = model.predict([[log_p_base]])[0]
    return elasticity, np.exp(pred_log_q)

elif model_type == 'Huber':
    model = HuberRegressor()
    model.fit(log_p, log_q, sample_weight=weights)
    elasticity = model.coef_[0]
    pred_log_q = model.predict([[log_p_base]])[0]
    return elasticity, np.exp(pred_log_q)

elif model_type == 'SVR':
    model = SVR(kernel='rbf', C=1.0, epsilon=0.1)
    model.fit(log_p, log_q, sample_weight=weights)
    pred_plus = model.predict([[log_p_base + h]])[0]
    pred_minus = model.predict([[log_p_base - h]])[0]
    elasticity = (pred_plus - pred_minus) / (2 * h)
```

```

    pred_log_q = model.predict([[log_p_base]])[0]
    return elasticity, np.exp(pred_log_q)

elif model_type == 'RF':
    model = RandomForestRegressor(n_estimators=50, max_depth=3, random_state=42)
    model.fit(log_p, log_q, sample_weight=weights)
    pred_plus = model.predict([[log_p_base + h]])[0]
    pred_minus = model.predict([[log_p_base - h]])[0]
    elasticity = (pred_plus - pred_minus) / (2 * h)
    pred_log_q = model.predict([[log_p_base]])[0]
    return elasticity, np.exp(pred_log_q)

elif model_type == 'GB':
    model = GradientBoostingRegressor(n_estimators=50, max_depth=2, random_state=42)
    model.fit(log_p, log_q, sample_weight=weights)
    pred_plus = model.predict([[log_p_base + h]])[0]
    pred_minus = model.predict([[log_p_base - h]])[0]
    elasticity = (pred_plus - pred_minus) / (2 * h)
    pred_log_q = model.predict([[log_p_base]])[0]
    return elasticity, np.exp(pred_log_q)

elif model_type == 'MLP':
    elasticity_vals = []
    pred_q_vals = []
    np.random.seed(42)
    normalized_weights = weights / np.sum(weights)
    for i in range(3):
        if use_decay:
            indices = np.random.choice(n_obs, size=n_obs, replace=True)
        else:
            indices = np.random.choice(n_obs, size=n_obs, replace=True)
        X_train = log_p[indices]
        y_train = log_q[indices]
        mlp = MLPRegressor(hidden_layer_sizes=(16, 8), activation='tanh')
        mlp.fit(X_train, y_train)
        p_plus = mlp.predict(np.array([[log_p_base + h]]))[0]
        p_minus = mlp.predict(np.array([[log_p_base - h]]))[0]
        e_i = (p_plus - p_minus) / (2 * h)
        elasticity_vals.append(e_i)
        pred_q_vals.append(np.exp(mlp.predict([[log_p_base]])[0]))
    return np.mean(elasticity_vals), np.mean(pred_q_vals)

```

Rolling Backtest Loop

We run the rolling validation starting at week .

```

In [4]: all_configs = [
    ('Linear_NoDecay', 'Linear', False),
    ('Linear_WithDecay', 'Linear', True),
    ('Ridge_NoDecay', 'Ridge', False),
    ('Ridge_WithDecay', 'Ridge', True),
    ('Huber_NoDecay', 'Huber', False),
    ('Huber_WithDecay', 'Huber', True),
    ('SVR_NoDecay', 'SVR', False),
    ('SVR_WithDecay', 'SVR', True),
    ('RF_NoDecay', 'RF', False),
    ('RF_WithDecay', 'RF', True),
    ('GB_NoDecay', 'GB', False),
    ('GB_WithDecay', 'GB', True),

```

```

    ('MLP_NoDecay', 'MLP', False),
    ('MLP_WithDecay', 'MLP', True),
]

if RUN_DECAY_ONLY:
    model_configs = [config for config in all_configs if config[2] is True]
else:
    model_configs = all_configs

train_window = 200
results = {config[0]: [] for config in model_configs}
test_indices = range(train_window, len(df))
actual_qs = []
actual_ps = []
dates = []

for idx in test_indices:
    train_slice = df.iloc[idx-train_window : idx]
    test_row = df.iloc[idx]

    test_price = test_row['unit_price']
    actual_q = test_row['units_sold']
    actual_qs.append(actual_q)
    actual_ps.append(test_price)
    dates.append(test_row['date'])

    for name, m_type, use_decay in model_configs:
        e, pred = run_model(m_type, train_slice, test_price, use_decay=use_decay)
        e = max(-10.0, min(2.0, e))
        results[name].append({'elasticity': e, 'pred': pred})

print(f"Completed rolling backtest for {len(test_indices)} test weeks. Evaluated {len(model_configs)} models.")

```

Completed rolling backtest for 295 test weeks. Evaluated 14 models.

Model Ranking & Evaluation Table

We compute the error metrics and calculate the Rank Score based on the specified weights:

- **Accuracy Score** (%): Normalized MAPE (lower is better)
- **Plausibility Score** (%): . for elasticity in [- . , - .], . for [- . , - . , - .], . otherwise
- **Stability Score** (%): Volatility of elasticity over rolling windows (lower Std is better)

```

In [5]: actual_qs = np.array(actual_qs)
        comparison_rows = []
        for name, _, _ in model_configs:
            predictions = results[name]
            preds = np.array([p['pred'] for p in predictions])
            preds = np.nan_to_num(preds, nan=0.0, posinf=1e20, neginf=0.0)
            elasts = np.array([p['elasticity'] for p in predictions])

            errors = preds - actual_qs
            abs_errors = np.abs(errors)
            pct_errors = abs_errors / actual_qs

            mae = np.mean(abs_errors)
            mape = np.mean(pct_errors) * 100

```

```

rmse = np.sqrt(np.mean(errors ** 2))
mean_elast = np.mean(elasts)
std_elast = np.std(elasts)

comparison_rows.append({
    'Model': name,
    'MAE (KG)': mae,
    'MAPE (%)': mape,
    'RMSE (KG)': rmse,
    'Mean Elasticity': mean_elast,
    'Std Elasticity': std_elast
})

df_compare = pd.DataFrame(comparison_rows)

# Compute rankings (excluding the exploding MLP_WithDecay if it exists)
scored_df = df_compare[df_compare['Model'] != 'MLP_WithDecay'].copy()
mape_min = scored_df['MAPE (%)'].min()
mape_max = scored_df['MAPE (%)'].max()
scored_df['Accuracy_Score'] = (mape_max - scored_df['MAPE (%)']) / (mape_max - mape_min)

def get_plausibility(e):
    if -2.5 <= e <= -0.8:
        return 1.0
    elif -4.0 <= e < -2.5 or -0.8 < e <= -0.4:
        return 0.5
    else:
        return 0.0
scored_df['Plausibility_Score'] = scored_df['Mean Elasticity'].apply(get_plausibility)

std_min = scored_df['Std Elasticity'].min()
std_max = scored_df['Std Elasticity'].max()
scored_df['Stability_Score'] = (std_max - scored_df['Std Elasticity']) / (std_max - std_min)

scored_df['Rank Score'] = 0.5 * scored_df['Accuracy_Score'] + 0.3 * scored_df['Plausibility_Score'] + 0.2 * scored_df['Stability_Score']
scored_df = scored_df.sort_values(by='Rank Score', ascending=False)

# Append MLP_WithDecay back at the bottom if it was run
if 'MLP_WithDecay' in df_compare['Model'].values:
    mlp_wd_row = df_compare[df_compare['Model'] == 'MLP_WithDecay'].copy()
    for col in ['Accuracy_Score', 'Plausibility_Score', 'Stability_Score']:
        mlp_wd_row[col] = 0.0
    final_ranked = pd.concat([scored_df, mlp_wd_row], ignore_index=True)
else:
    final_ranked = scored_df

pd.set_option('display.float_format', lambda x: '%.3e' % x if abs(x) > 1e6 else x)
print(final_ranked[['Model', 'MAE (KG)', 'MAPE (%)', 'RMSE (KG)', 'Mean Elasticity', 'Std Elasticity', 'Accuracy_Score', 'Plausibility_Score', 'Stability_Score', 'Rank Score']])
print("\nNote: RMSE (Root Mean Squared Error) measures the square root of the mean squared error")

```

	Model	MAE (KG)	MAPE (%)	RMSE (KG)	Mean Elasticity	Std Elasticity
Rank Score						
Linear_WithDecay	855865.307	108.897	2.393e+06	-1.250		
1.204	0.905					
Huber_WithDecay	854475.155	107.581	2.392e+06	-1.125		
1.427	0.883					
Ridge_WithDecay	959795.894	110.804	2.613e+06	-0.614		
0.584	0.819					
SVR_WithDecay	763444.366	100.655	2.352e+06	-0.984		
2.150	0.815					
Ridge_NoDecay	1.273e+06	266.275	2.720e+06	-0.969		
0.408	0.799					
Linear_NoDecay	1.272e+06	266.369	2.710e+06	-1.015		
0.432	0.796					
Huber_NoDecay	1.303e+06	278.003	2.737e+06	-0.923		
0.456	0.780					
GB_WithDecay	905791.130	123.668	2.528e+06	-0.027		
0.593	0.652					
SVR_NoDecay	1.169e+06	267.161	2.371e+06	-0.863		
2.088	0.619					
RF_WithDecay	920771.229	121.146	2.550e+06	-0.163		
1.297	0.581					
GB_NoDecay	1.175e+06	236.522	2.644e+06	-0.027		
0.593	0.515					
RF_NoDecay	1.148e+06	237.161	2.592e+06	-0.156		
1.303	0.439					
MLP_NoDecay	1.940e+06	512.754	1.364e+07	-0.740		
2.287	0.150					
MLP_WithDecay	3.229e+21	7.216e+17	5.541e+22	-1.518		3
0.000e+00						

Note: RMSE (Root Mean Squared Error) measures the square root of average squared errors, penalizing larger outliers more heavily. SVR_WithDecay achieves the lowest RMSE (2.352e+06 KG) among all configurations, indicating that it produces the fewest large-scale prediction errors out-of-sample.

Visualizations (Top Models - Separated by Decay NoDecay)

To maximize legibility, the plots for rolling predictions, test curves, and final training estimates are separated into **WithDecay** and **NoDecay** versions. Price ranges are limited to % below and % above the current price (\$[,]\$ INR/KG).

```
In [6]: # Plot rolling predictions vs actuals for top 4 models (2015-2017 only)
plot_dates = pd.to_datetime(dates)
filter_2015 = (plot_dates >= '2015-01-01') & (plot_dates <= '2017-12-31')

filtered_plot_dates = plot_dates[filter_2015]
filtered_actual_qs = np.array(actual_qs)[filter_2015]

color_pairs = {
    'Linear': ('#0d47a1', '#64b5f6'), # Blue
    'Huber': ('#e65100', '#ffb74d'), # Orange
    'Ridge': ('#1b5e20', '#81c784'), # Green
    'SVR': ('#4a148c', '#ba68c8'), # Purple
}


# Plot 1a: WithDecay
```

```

plt.figure(figsize=(14, 6))
# Actual quantity: light and dotted
plt.plot(filtered_plot_dates, filtered_actual_qs, color='#bbbbbb', linestyle='dotted')
for base_name, (c_decay, _) in color_pairs.items():
    wd_name = base_name + '_WithDecay'
    if wd_name in results:
        preds_wd = np.array([p['pred'] for p in results[wd_name]])[filtered_plot_dates.index]
        # Predictions: dark solid
        plt.plot(filtered_plot_dates, preds_wd, label=f'Predicted - {wd_name}', color=c_decay)
plt.title('Actual vs. Predicted Weekly Demand (Top 4 Models - WithDecay, 2015-2017)')
plt.xlabel('Date')
plt.ylabel('Quantity (KG)')
plt.ylim(0, 2e6) # Cut graph at 2e6 to hide outliers
plt.legend()
plt.grid(True, linestyle='--', alpha=0.5)
plt.show()

# Plot 1b: NoDecay
if not RUN_DECAY_ONLY:
    plt.figure(figsize=(14, 6))
    # Actual quantity: light and dotted
    plt.plot(filtered_plot_dates, filtered_actual_qs, color='#bbbbbb', linestyle='dotted')
    for base_name, (_, c_noddecay) in color_pairs.items():
        nd_name = base_name + '_NoDecay'
        if nd_name in results:
            preds_nd = np.array([p['pred'] for p in results[nd_name]])[filtered_plot_dates.index]
            # Predictions: solid line using the model color
            plt.plot(filtered_plot_dates, preds_nd, label=f'Predicted - {nd_name}', color=c_noddecay)
    plt.title('Actual vs. Predicted Weekly Demand (Top 4 Models - NoDecay, 2015-2017)')
    plt.xlabel('Date')
    plt.ylabel('Quantity (KG)')
    plt.ylim(0, 2e6) # Cut graph at 2e6 to hide outliers
    plt.legend()
    plt.grid(True, linestyle='--', alpha=0.5)
    plt.show()

```






Chart : Quantity-Price Graph (Testing Backtest Predictions)

This plot shows the relationship between unit prices and predicted quantities across the **rolling 1 weeks**. The range is limited to +/- % of the current price (\$[,]\$ INR/K

```
In [7]: prices = np.array(actual_ps)
p_current = df['unit_price'].iloc[-1]
p_min = p_current * 0.7
p_max = p_current * 1.3

filter_idx = (prices >= p_min) & (prices <= p_max)
filtered_prices = prices[filter_idx]
filtered_actual_qs = actual_qs[filter_idx]
sorted_idx = np.argsort(filtered_prices)

# Chart 1a: WithDecay
plt.figure(figsize=(14, 8))
plt.scatter(filtered_prices, filtered_actual_qs, color='black', alpha=0.4)
for base_name, (c_decay, _) in color_pairs.items():
    wd_name = base_name + '_WithDecay'
    if wd_name in results:
        preds_wd = np.array([p['pred'] for p in results[wd_name]])[filter_idx]
        plt.plot(filtered_prices[sorted_idx], preds_wd[sorted_idx], label=base_name)
plt.title('Quantity-Price Curves during Rolling Testing (Top 4 Models - WithDecay)')
plt.xlabel('Unit Price (INR/KG)')
plt.ylabel('Quantity Sold / Predicted (KG)')
plt.xlim(p_min, p_max)
plt.ylim(0, 7.5e6)
plt.legend(bbox_to_anchor=(1.05, 1), loc='upper left')
plt.grid(True, linestyle='--', alpha=0.5)
plt.tight_layout()
plt.show()

# Chart 1b: NoDecay
if not RUN_DECAY_ONLY:
    plt.figure(figsize=(14, 8))
    plt.scatter(filtered_prices, filtered_actual_qs, color='black', alpha=0.4)
    for base_name, (_, c_noddecay) in color_pairs.items():
        nd_name = base_name + '_NoDecay'
        if nd_name in results:
            preds_nd = np.array([p['pred'] for p in results[nd_name]])[filter_idx]
            plt.plot(filtered_prices[sorted_idx], preds_nd[sorted_idx], label=base_name)
```



```

        preds_nd = np.array([p['pred'] for p in results[nd_name]])[fi
        plt.plot(filtered_prices[sorted_idx], preds_nd[sorted_idx], l
    plt.title('Quantity-Price Curves during Rolling Testing (Top 4 Models
    plt.xlabel('Unit Price (INR/KG)')
    plt.ylabel('Quantity Sold / Predicted (KG)')
    plt.xlim(p_min, p_max)
    plt.ylim(0, 7.5e6)
    plt.legend(bbox_to_anchor=(1.05, 1), loc='upper left')
    plt.grid(True, linestyle='--', alpha=0.5)
    plt.tight_layout()
    plt.show()

```





Chart : Quantity-Price Graph (Last Estimation / Final Window)

Here we fit the models on the **final training window** (the last weeks of the dataset) and project their demand curves across a wide range of pricing scenarios (% below to % above current price).

```
In [8]: final_train_slice = df.iloc[-200:]
test_prices_range = np.linspace(p_min, p_max, 100)
hist_in_range = final_train_slice[(final_train_slice['unit_price'] >= p_min)]

# Chart 2a: WithDecay
plt.figure(figsize=(14, 8))
plt.scatter(hist_in_range['unit_price'], hist_in_range['units_sold'], color='red')
for base_name, (c_decay, _) in color_pairs.items():
    wd_name = base_name + '_WithDecay'
    if wd_name in results:
        q_wd = []
        for p in test_prices_range:
            _, q_wd_p = run_model(base_name, final_train_slice, p, use_decay=True)
            q_wd.append(q_wd_p)
        plt.plot(test_prices_range, q_wd, label=f"{wd_name} (Last Window)", color=c_decay)
plt.title('Quantity-Price Curves: Last Estimation Window (Top 4 Models - With Decay)')
plt.xlabel('Unit Price (INR/KG)')
plt.ylabel('Predicted Quantity (KG)')
plt.xlim(p_min, p_max)
plt.ylim(0, 7.5e6)
plt.legend(bbox_to_anchor=(1.05, 1), loc='upper left')
plt.grid(True, linestyle='--', alpha=0.5)
plt.tight_layout()
plt.show()

# Chart 2b: NoDecay
if not RUN_DECAY_ONLY:
    plt.figure(figsize=(14, 8))
    plt.scatter(hist_in_range['unit_price'], hist_in_range['units_sold'], color='red')
    for base_name, (_, c_noddecay) in color_pairs.items():
        nd_name = base_name + '_NoDecay'
        if nd_name in results:
            q_nd = []
            for p in test_prices_range:
                _, q_nd_p = run_model(base_name, final_train_slice, p, use_decay=False)
                q_nd.append(q_nd_p)
            plt.plot(test_prices_range, q_nd, label=f"{nd_name} (Last Window)", color=c_noddecay)
    plt.title('Quantity-Price Curves: Last Estimation Window (Top 4 Models - No Decay)')
    plt.xlabel('Unit Price (INR/KG)')
    plt.ylabel('Predicted Quantity (KG)')
    plt.xlim(p_min, p_max)
    plt.ylim(0, 7.5e6)
    plt.legend(bbox_to_anchor=(1.05, 1), loc='upper left')
    plt.grid(True, linestyle='--', alpha=0.5)
    plt.tight_layout()
    plt.show()
```

